

Customer Intelligence and Big Data

Alan Rijnders and Lorenzo Severi

11/4/2021

1. Introduction

We, Alan Rijnders and Lorenzo Severi have been appointed from Alpha Bank to help tackling one of the biggest problems physical banks are facing now, being the significant number of customers shifting towards online financial institutions, such as N26 and illimity Bank.

If Alpha Bank aims to survive in the financial market in the long run, a proactive way of predicting the customers willing to leave their service is essential. In fact, by discovering beforehand those customers, Alpha Bank will be in the position to find solutions aimed at better suiting their needs by eventually persuading them to remain.

One might argue that some scholars have recently suggested companies avoiding focusing just on retaining customers; long-term customers seem indeed to be extremely costly, as they tend to expect lower prices in exchange for their loyalty (Anderson & Jap, 2005). Nonetheless, it is common knowledge that, if feasible, companies should find ways to persuade them to remain, but without forgetting to keep an eye on their profitability.

2. Methodology

The database provided consist of labeled data, meaning that it represents a case of supervised learning. As we will be dealing with a prediction of a binary category (i.e., “0 à remain” or “1 à leave”), the problem involves classification, as the outcome variable will not be continuous. During our classes, we have seen several statistical methods suitable for this kind of data, namely Logistic Regression, Decision Tree and Random Forest.

Logistic regression is a predictive analysis used to explain the relationship between one non-metric dependent variable and one (or several) nominal, ordinal, interval or ratio independent variables. Instead of predicting the dependent variable, this technique predicts the probability of the dependent variable to be true (Statistics Solutions, n.d.; Edgar & Manz, 2017).

Decision trees (also known as CART trees) are graphic representation of different alternative solutions that are suitable to tackle a problem. It is usually used to determine the most effective courses of action. All decision trees start with a single node that is connected through branches into possible outcomes. All the different outcomes lead to additional nodes which branch off into other possibilities (OmniSci, n.d.; Lucidchart, n.d.).

Random forest is a collection of decision trees trained with a bootstrapped technique, so that the ensemble boosts the accuracy of the decision trees outcome (Donges, 2021).

3. Data preparation

To be able to analyze the dataset, we first need to set the working directory where the file is located, and subsequently to open the .csv file. Then, we need to install and call in several free libraries, as they are required for conducting our analysis. With regards to the library “caret”, it is important to install it by means of the dependencies, as we need them in order to calculate sensitivity and specificity.

```
#read in data
data <- read.csv("ch.csv")
library(dplyr)
```

```

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(ggplot2)
#install.packages('caret', dependencies = TRUE)
#install with dependencies = TRUE is important for the calculation of sensitivity and specificity
library(caret)

## Loading required package: lattice
library(corrplot)

## corrplot 0.90 loaded
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --

## v tibble  3.1.4      v purrr   0.3.4
## v tidyr   1.1.4      v stringr 1.4.0
## v readr   2.0.2      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## x purrr::lift()    masks caret::lift()

library(repr)
library(caTools)
library(pROC)

## Type 'citation("pROC")' for a citation.

##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##   cov, smooth, var

library(rpart)
library(rpart.plot)
library(ggpubr)
library(randomForest)

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##

```

```
##      margin
## The following object is masked from 'package:dplyr':
##
##      combine
```

By briefly looking at the variables, we see how some need to be re-recorded. What is more, we considered the first two columns of our database (i.e., “X” and “CLIENTUM”) as neglectable and thus we decided to crop them to obtain a cleaner dataset. Then, we proceed with the transformation of several variables (i.e., “Gender”, “Education_Level”, “Marital_Status”, “Income_Category” and “Card_Category”) into factor variables, so that we will be able to analyze them accurately. The last step is to record the variable “Attrition_Flag”, which contains the two levels “Existing Customer” and “Attrited Customer”, into a binary variable, with the value of “0” and “1” respectively. As there seems to be no possible detrimental errors in the dataset (e.g., non-response error, out-of-range error, ...), we will be further proceeding by assuming that the quantitative inferences we will take on the outcome will not be unvalidated.

```
#we drop the X and client number column
data <- subset(data, select=-c(X,CLIENTNUM))
```

```
#summary of the dataset
summary(data)
```

```
## Attrition_Flag      Customer_Age      Gender      Dependent_count
## Length:10127      Min.      :26.00      Length:10127      Min.      :0.000
## Class :character      1st Qu.:41.00      Class :character      1st Qu.:1.000
## Mode  :character      Median :46.00      Mode  :character      Median :2.000
##                               Mean  :46.33              Mean  :2.346
##                               3rd Qu.:52.00              3rd Qu.:3.000
##                               Max.   :73.00              Max.   :5.000
## Education_Level      Marital_Status      Income_Category      Card_Category
## Length:10127      Length:10127      Length:10127      Length:10127
## Class :character      Class :character      Class :character      Class :character
## Mode  :character      Mode  :character      Mode  :character      Mode  :character
##
##
##
## Months_on_book      Total_Relationship_Count      Months_Inactive_12_mon
## Min.      :13.00      Min.      :1.000              Min.      :0.000
## 1st Qu.:31.00      1st Qu.:3.000              1st Qu.:2.000
## Median :36.00      Median :4.000              Median :2.000
## Mean   :35.93      Mean   :3.813              Mean   :2.341
## 3rd Qu.:40.00      3rd Qu.:5.000              3rd Qu.:3.000
## Max.   :56.00      Max.   :6.000              Max.   :6.000
## Contacts_Count_12_mon      Credit_Limit      Total_Trans_Amt      Total_Trans_Ct
## Min.      :0.000              Min.      : 1438      Min.      : 510      Min.      : 10.00
## 1st Qu.:2.000              1st Qu.: 2555      1st Qu.: 2156      1st Qu.: 45.00
## Median :2.000              Median : 4549      Median : 3899      Median : 67.00
## Mean   :2.455              Mean   : 8632      Mean   : 4404      Mean   : 64.86
## 3rd Qu.:3.000              3rd Qu.:11068      3rd Qu.: 4741      3rd Qu.: 81.00
## Max.   :6.000              Max.   :34516      Max.   :18484      Max.   :139.00
## Avg_Utilization_Ratio
## Min.      :0.0000
## 1st Qu.:0.0230
## Median :0.1760
## Mean   :0.2749
## 3rd Qu.:0.5030
```

```
## Max. :0.9990
```

Data exploration

The first step of data exploration consists of reading and visualizing the data: there were 10,127 observations with 18 variables, ranging from non-metric to metric ones.

With regards to the socio-demographic characteristics, the average age was forty-six, with a minimum age of twenty-six and a maximum age of seventy-three, and the database consists of 52.9% female (5358) and 47.1% males (4769). With regards to the number of dependents, the average is slightly more than two people, with a maximum of five. 70.3% of the customers obtained at least an high-school diploma, but a significant 14.7% is not educated, and 15% have an unknown status. Concerning the marital status, 46.3% of people are married, 38.9% are single and the remaining 14.8% are either divorced or with an unknown status. The income category is considerable diversified: 35.2% has an annual income that is less than \$40k, 17.7% \$40k - \$60k, 13.8% \$60k - \$80k, 15.2% \$80 \$120k, 7.2% more than \$120k, and 11% have an unknown income status.

Concerning the variables related to the product, the data depicts that more than 93% of the customers have blue card, the average time of bank-customer relationship is 35 months with an average of almost four active products. In the graphs below, we show the different distribution of the aforementioned variables by means of their status with the Alpha Bank (i.e., “Attrited customer” vs “Existing customer”).

```
table(data$Attrition_Flag)
```


```
##
## Attrited Customer Existing Customer
##                1627                8500
```

We see that our dependent variable shows that out of the 10,127 customers, 83.9% (8,500 customers) have remained, whereas 16.1% (1,627 customers) have left.

Similarly we transform the variables Gender, Education Level, Marital Status, Income Category and Card Category to factor variables.

```
data <- transform(
  data,
  Attrition_Flag = as.factor(Attrition_Flag),
  Gender = as.factor(Gender),
  Education_Level = as.factor(Education_Level),
  Marital_Status = as.factor(Marital_Status),
  Income_Category = as.factor(Income_Category),
  Card_Category = as.factor(Card_Category))
```


```
nv <- sapply(data, is.numeric)
cormat <- cor(data[,nv])
corrplot::corrplot(cormat, title = "Correlation of Numeric Variables")
```



With regards to the correlation matrix, we see in Figure “Correlation of Numeric Variables” how some variables seem to be highly correlated. However, that was highly expected due to the fact that having a high number of variables increases the likelihood of finding overlaps among them. The variables involved are: “Months_on_book” with “Customer_Age” (positively correlated), “Total_Trans_Amt” with “Total_Trans_Ct” (positively correlated) and “Avg_Utilization_Ratio” with “Credit_Limit” (negatively correlated). Since they do not measure the same aspects, we proceed by considering all of them as adding significant information into our model, and thus we are not excluding one of them.

In the graphs below, we show the different distribution of the aforementioned variables by means of their status with the Alpha Bank (i.e., “Attrited customer” vs “Existing customer”).


```
fig1 <- ggarrange(ggplot(data, aes(x=Gender,fill=Attrition_Flag))+ geom_bar() ,
  ggplot(data, aes(x=Marital_Status,fill=Attrition_Flag))+ geom_bar(position = 'fill'))
print(fig1)
```



```


annotate_figure(fig1, bottom = text_grob("Attrition Percentage in Gender, Marital Status and Card Category",

```



Attrition Percentage in Gender, Marital Status and Card Category


```
fig2 <- ggplot(data, aes(x=Education_Level,fill=Attrition_Flag))+ geom_bar(position = 'fill')
print(fig2)
```



```


annotate_figure(fig2, bottom = text_grob("Attrition Percentage for different levels of education", col

```

Attrition Percentage for different levels of education


```
fig3 <- ggplot(data, aes(x=Income_Category, fill=Attrition_Flag)) + geom_bar(position = 'fill')
print(fig3)
```



```


annotate_figure(fig3, bottom = text_grob("Attrition Percentage for different income levels", col = "black",

```

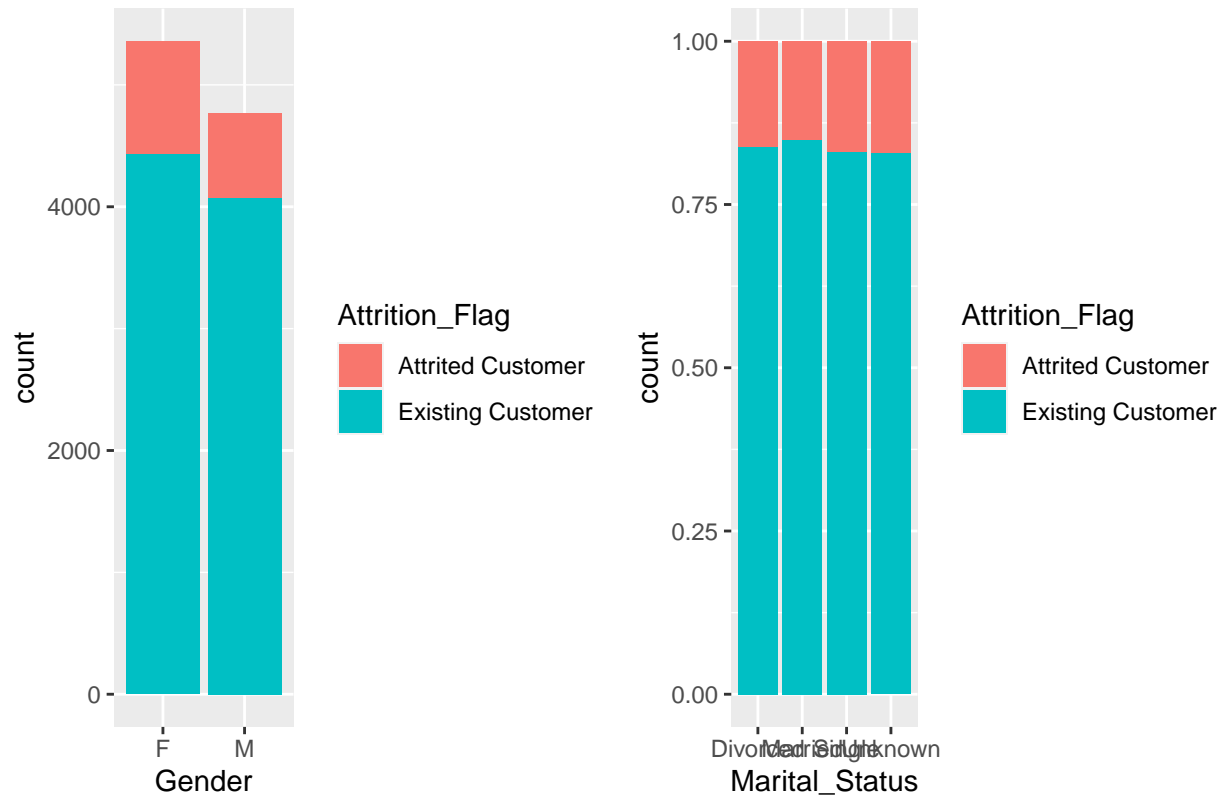


Attrition Percentage for different income levels

```
fig4 <- ggplot(data, aes(x=Card_Category,fill=Attrition_Flag))+ geom_bar(position = 'fill')
print(fig4)
```

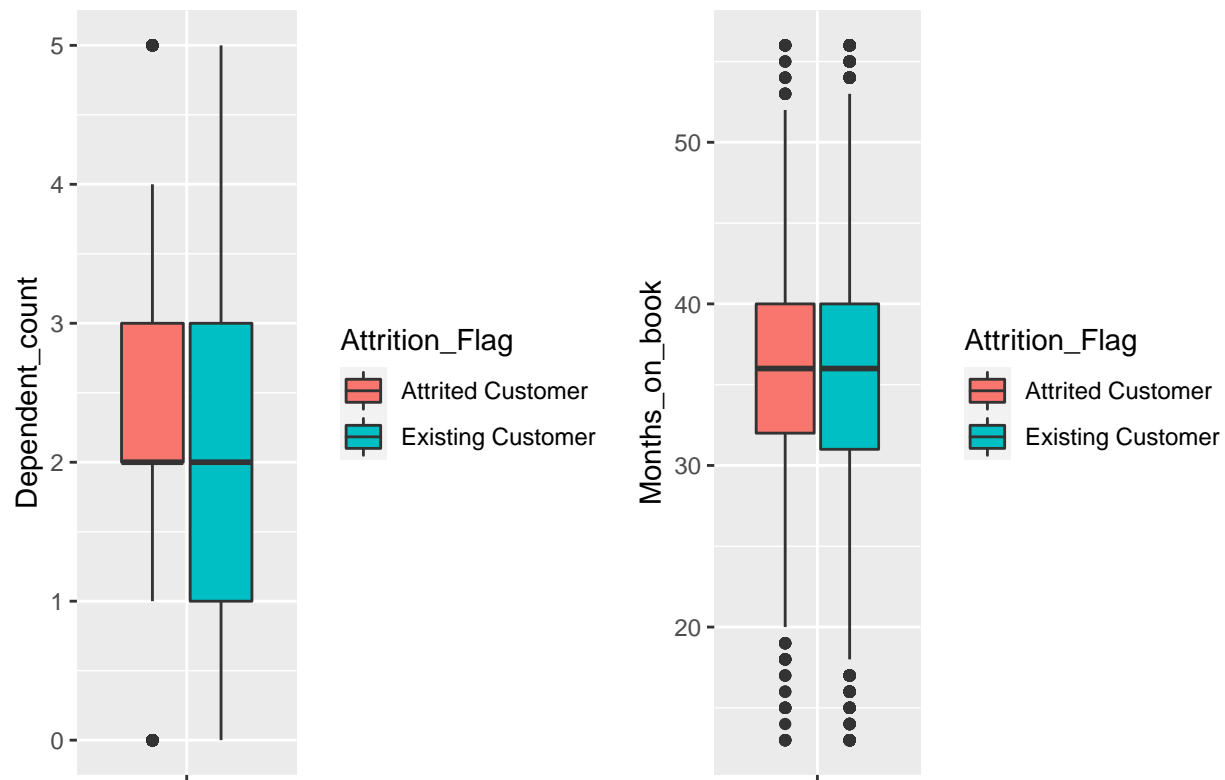


```
annotate_figure(fig1, bottom = text_grob("Attrition Percentage for different cardholder categories", c
```



Attrition Percentage for different cardholder categories

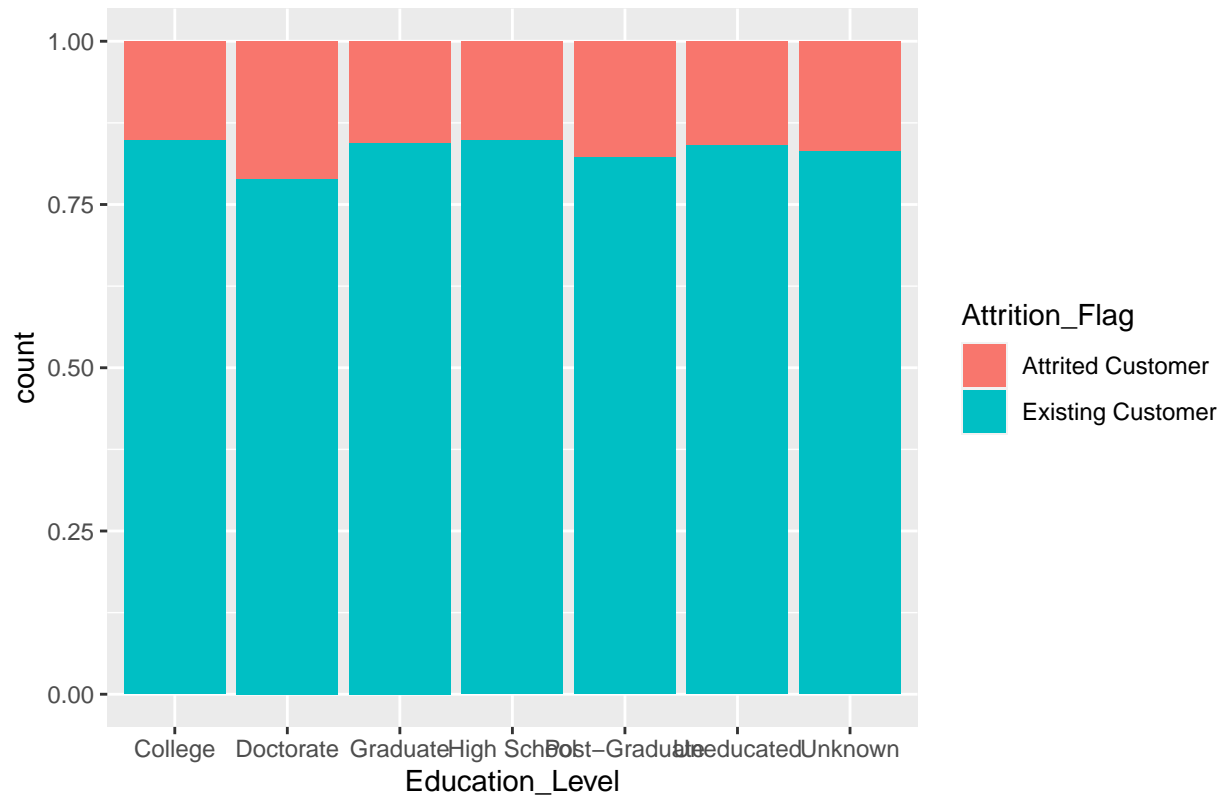
```
fig5 <- ggarrange(
  ggplot(data, aes(y= Dependent_count, x = "", fill = Attrition_Flag))
+geom_boxplot() + xlab(" "),
  ggplot(data, aes(y= Months_on_book, x = "", fill = Attrition_Flag))
+geom_boxplot() + xlab(" "))
print(fig5)
```



```

annotate_figure(fig2, bottom = text_grob("Attrition Percentage for different number of dependents in f

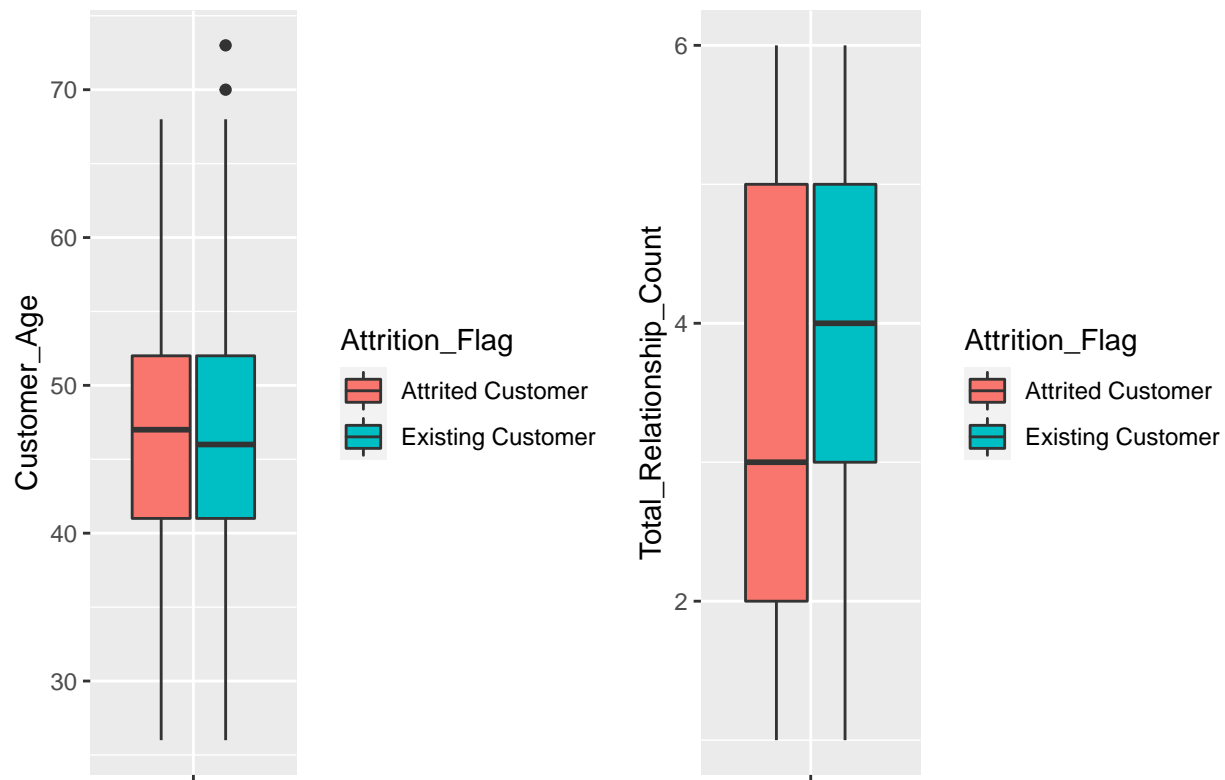
```



ge for different number of dependents in family and different number o


```
fig6 <- ggarrange(
  ggplot(data, aes(y= Customer_Age, x = "", fill = Attrition_Flag))
+geom_boxplot() + xlab(" "),
  ggplot(data, aes(y= Total_Relationship_Count, x = "", fill = Attrition_Flag))
+geom_boxplot() + xlab(" "))

print(fig6)
```



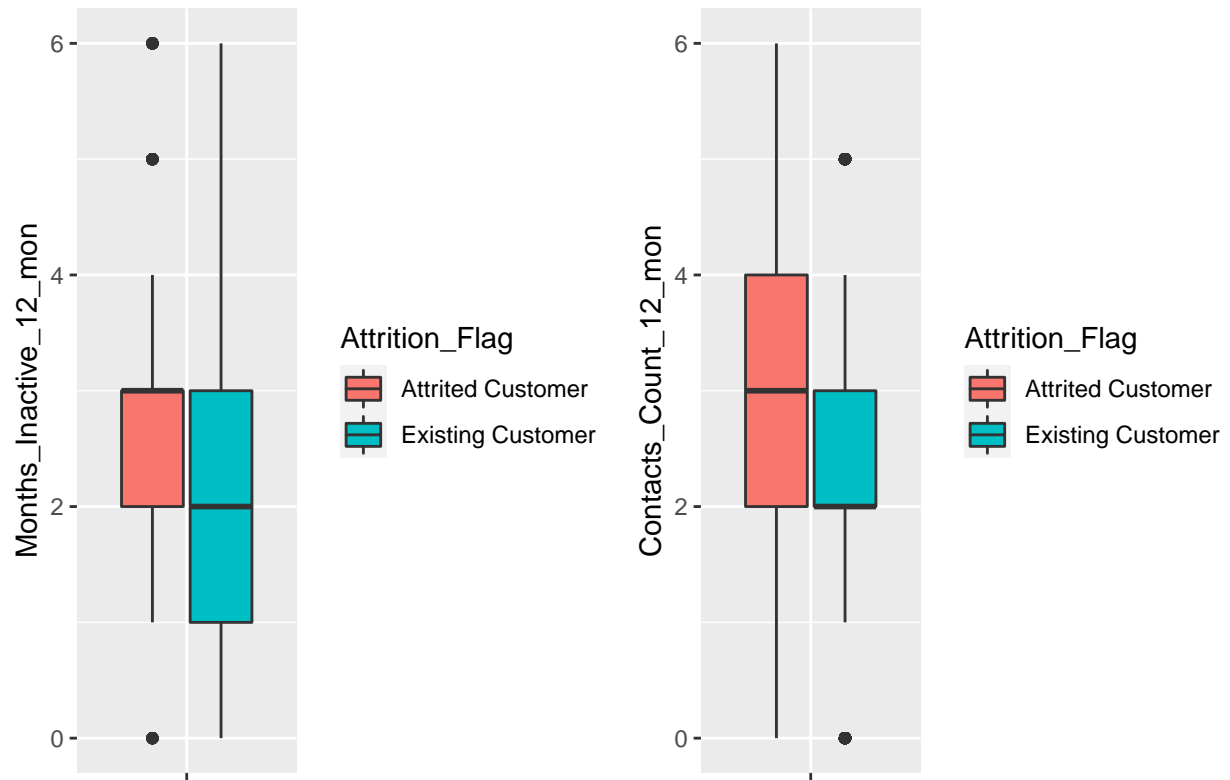
```

annotate_figure(fig6, bottom = text_grob("Attrition Percentage in Age and Relationship counts", col = "green", size = 12, weight = "bold"),
  
```





Attrition Percentage in Age and Relationship counts

```
fig7 <- ggarrange(
  ggplot(data, aes(y= Months_Inactive_12_mon, x = "", fill = Attrition_Flag))
+geom_boxplot() + xlab(" "),
  ggplot(data, aes(y= Contacts_Count_12_mon, x = "", fill = Attrition_Flag))
+geom_boxplot() + xlab(" "))
print(fig7)
```




```
annotate_figure(fig7, bottom = text_grob("Attrition Percentage in inactivity and number of contracts",
```



Attrition Percentage in inactivity and number of contracts

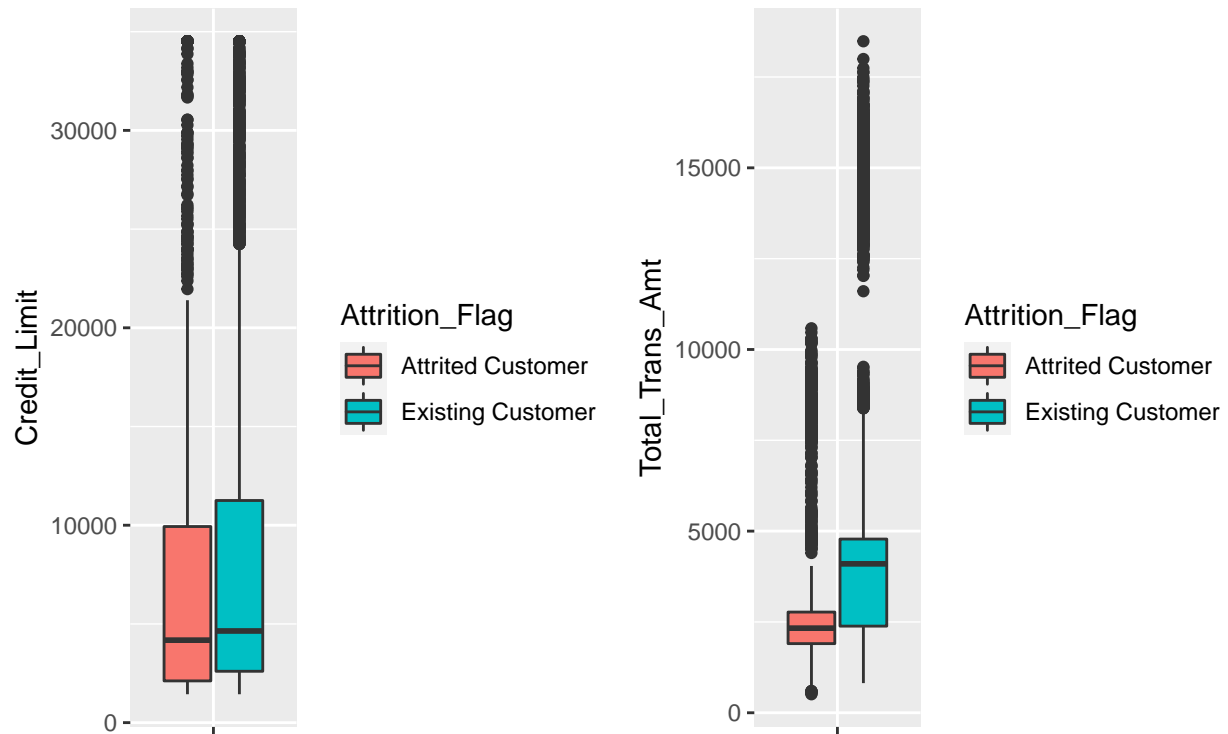
```
fig8 <- ggarrange(
  ggplot(data, aes(y= Credit_Limit, x = "", fill = Attrition_Flag))
+geom_boxplot() + xlab(" "),
  ggplot(data, aes(y= Total_Trans_Amt, x = "", fill = Attrition_Flag))
+geom_boxplot() + xlab(" "))
print(fig8)
```



```

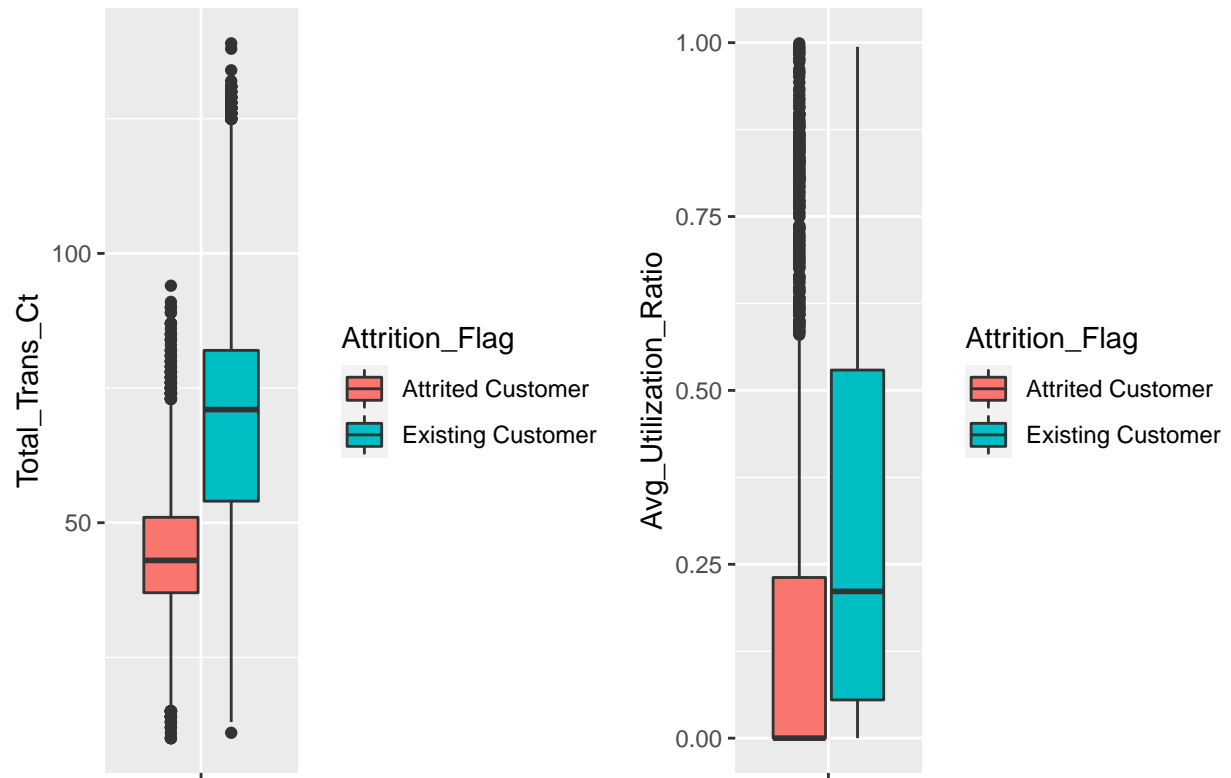
annotate_figure(fig8, bottom = text_grob("Attrition Percentage in different levels of credit limit tran

```

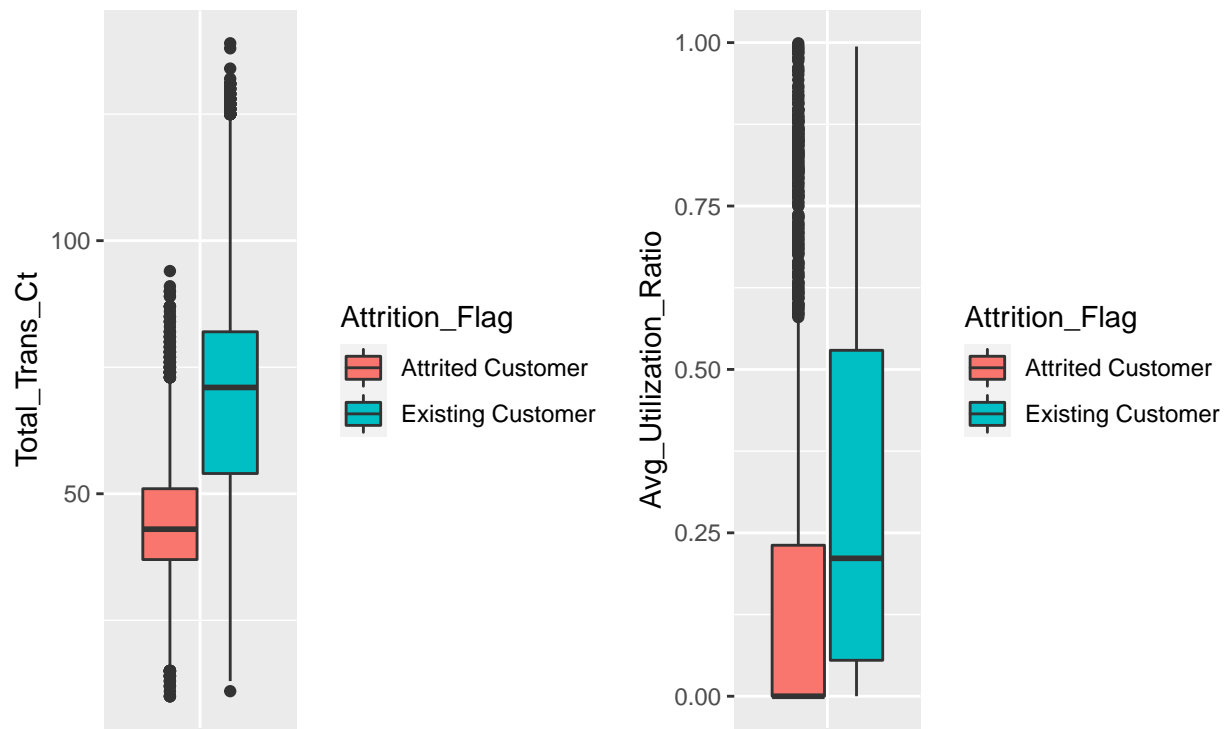


Attrition Percentage in different levels of credit limit transaction levels

```
fig9 <- ggarrange(
  ggplot(data, aes(y= Total_Trans_Ct, x = "", fill = Attrition_Flag))
+geom_boxplot() + xlab(" "),
  ggplot(data, aes(y= Avg_Utilization_Ratio, x = "", fill = Attrition_Flag))
+geom_boxplot() + xlab(" "))
print(fig9)
```



`annotate_figure(fig9, bottom = text_grob("Attrition Percentage in number of transactions and utilization"))`



Attrition Percentage in number of transactions and utilization ratio

```
##      Attrition_Flag      Customer_Age      Gender
##      "factor"           "integer"        "factor"
##      Dependent_count    Education_Level   Marital_Status
##      "integer"          "factor"         "factor"
##      Income_Category    Card_Category    Months_on_book
##      "factor"           "factor"         "integer"
## Total_Relationship_Count Months_Inactive_12_mon Contacts_Count_12_mon
## "integer"              "integer"        "integer"
##      Credit_Limit       Total_Trans_Amt   Total_Trans_Ct
##      "numeric"          "integer"         "integer"
##      Avg_Utilization_Ratio
##      "numeric"
```

It is essential to split the data into training and test sets. Indeed, we will first use the training data to build our model and then, when we are satisfied with the achievements, the remaining test data will be used to test our model, and thus to assess the power of our predictions. The training set is composed by 75% of the data, whereas the test set 25%.

```
smp_size <- floor(0.75 * nrow(data))

## set the seed to make your partition reproducible
set.seed(12345)
train_ind <- sample(seq_len(nrow(data)), size = smp_size)

train <- data[train_ind, ]
test  <- data[-train_ind, ]
```

```
prop.table(table(train$Attrition_Flag))
```

```
##
## Attrited Customer Existing Customer
##      0.158262      0.841738
```

Since Attrition_Flag is a character variable with two possible values: either “Existing Customer” or “Attrited Customer” for modeling purposes we recode this variable into a factor with levels 0 and 1, where 1 represents a customer that has left the company when the person is still a customer at the company.

```
#recoding attrition_flag
train$Attrition_Flag <- ifelse(train$Attrition_Flag=="Attrited Customer",1,0)
test$Attrition_Flag <- ifelse(test$Attrition_Flag=="Attrited Customer", 1,0)
```

We now can proceed with our first technique, namely Logistic regression. As we don't have any preference between the two errors (i.e., Type 1 or Type 2), we proceed by selecting the threshold of t to be 0.5, as it will predict the most likely outcome. Indeed, as explained at the beginning of the report, companies should be advised that spending too much money on retaining customers is often the wrong action to take (even though it is important to find ways to retain). Hence, we decided that false positive and false negative would have the same impact, as it would be costly to lose a customer if not considered as willing to go (false positive?), but at the same time it is costly to find ways to retain them so considering a customer that wants to go even though is not costly as well (I said he goes but he remained -> false negative).

```
#logistic regression
glm <- glm(Attrition_Flag ~., data = train, family = "binomial")
summary(glm)
```

```
##
## Call:
## glm(formula = Attrition_Flag ~ ., family = "binomial", data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.6806  -0.4127  -0.1936  -0.0800   3.3684
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      4.916e+00  4.913e-01  10.007 < 2e-16 ***
## Customer_Age     -5.264e-04  8.325e-03  -0.063  0.949579
## GenderM          -1.009e+00  1.652e-01  -6.105  1.03e-09 ***
## Dependent_count    1.094e-01  3.241e-02   3.375  0.000737 ***
## Education_LevelDoctorate  3.667e-01  2.200e-01   1.667  0.095536 .
## Education_LevelGraduate  1.094e-01  1.520e-01   0.720  0.471788
## Education_LevelHigh School -3.988e-02  1.620e-01  -0.246  0.805586
## Education_LevelPost-Graduate  2.768e-01  2.222e-01   1.246  0.212839
## Education_LevelUneducated  9.045e-02  1.703e-01   0.531  0.595397
## Education_LevelUnknown    1.984e-01  1.683e-01   1.179  0.238356
## Marital_StatusMarried    -5.234e-01  1.698e-01  -3.082  0.002054 **
## Marital_StatusSingle     1.005e-01  1.705e-01   0.589  0.555614
## Marital_StatusUnknown    1.198e-01  2.132e-01   0.562  0.574139
## Income_Category$40K - $60K -1.214e+00  2.208e-01  -5.499  3.82e-08 ***
## Income_Category$60K - $80K -9.105e-01  1.905e-01  -4.780  1.75e-06 ***
## Income_Category$80K - $120K -5.444e-01  1.776e-01  -3.065  0.002178 **
## Income_CategoryLess than $40K -1.048e+00  2.401e-01  -4.365  1.27e-05 ***
## Income_CategoryUnknown    -1.166e+00  2.517e-01  -4.634  3.58e-06 ***
## Card_CategoryGold         1.087e+00  3.964e-01   2.742  0.006111 **
```



```
## Card_CategoryPlatinum      2.162e+00  7.218e-01   2.996 0.002740 **
## Card_CategorySilver        5.485e-01  2.070e-01   2.650 0.008055 **
## Months_on_book             -1.004e-02  8.315e-03  -1.208 0.227239
## Total_Relationship_Count    -4.734e-01  3.007e-02 -15.746 < 2e-16 ***
## Months_Inactive_12_mon      4.953e-01  4.134e-02  11.982 < 2e-16 ***
## Contacts_Count_12_mon       5.425e-01  3.968e-02  13.672 < 2e-16 ***
## Credit_Limit                -6.085e-05  7.126e-06  -8.539 < 2e-16 ***
## Total_Trans_Amt             4.121e-04  2.497e-05  16.504 < 2e-16 ***
## Total_Trans_Ct              -1.139e-01  3.964e-03 -28.738 < 2e-16 ***
## Avg_Utilization_Ratio       -2.959e+00  1.844e-01 -16.048 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 6634.6  on 7594  degrees of freedom
## Residual deviance: 3958.3  on 7566  degrees of freedom
## AIC: 4016.3
##
## Number of Fisher Scoring iterations: 6
```

We first look at accuracy (i.e., number of corrected guesses), specificity (i.e., true negative rate), sensitivity (i.e., true positive rate), and precision on our training set, so that we can analyze how well the model was at predicting the outcome.

```
pred <- predict(glm, data = train, type = "response")
# confusion matrix on training set
conmat <- table(train$Attrition_Flag, pred >= 0.5)
#show confusion matrix
print(conmat)
```

```
##
##      FALSE TRUE
##  0   6149   244
##  1    590   612
```

```
accuracy <- (6149+612)/nrow(train)
specificity <- 6149/(6149+244)
sensitivity <- 612/(612+590)
precision <- 612/(612+244)
```

The accuracy seems to be pretty

Secondly we deploy the model on our test set to see how well our model predicts on new data. We again calculate the Accuracy, Specificity, Sensitivity and Precision of our model.

```
# observations on the test set
predtest <- predict(glm, newdata = test, type = "response")
conMattest <- table(test$Attrition_Flag, predtest >= 0.5)
#show confusion matrix
print(conMattest)
```

```
##
##      FALSE TRUE
##  0   2029    78
##  1    189   236
```

```

accuracytest <- (2029+236)/nrow(test)
specificitytest <- 2029/(2029+78)
sensitivitytest <- 236/(236+189)
precisiontest <- 236/(236+78)

```

Again, the accuracy seems to be pretty


We now plot the ROC curve of our model. The ROC curve measures the .. the larger the area the more accurate the model, with an area of 1 signifying a perfect model and an area of 0.5 or lower as a situation in which the model does not perform better than a random prediction.

```

par(mai=c(.9,.8,.2,.2))
plot(roc(test$Attrition_Flag, predtest), print.auc=TRUE,
     col="black", lwd=1, main="ROC curve", xlab="Specificity: true negative rate", ylab="Sensitivity: true positive rate")

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

```



The area under the curve equals to ... this is a very high value and is an indication that our model is performing quite well.

Next we calculate the variable importance for all of the variables we include in our model. The variable importance measures which variables are the main predictors of the loyalty of our customers. Insight into the main predictors could help give the company focus in developing marketing strategies to prevent attrition.

```


logisticvariableimportance <- varImp(glm, scale = FALSE)
print(logisticvariableimportance)

```

##	Overall
## Customer_Age	0.0632359
## GenderM	6.1052300
## Dependent_count	3.3753243
## Education_LevelDoctorate	1.6668892
## Education_LevelGraduate	0.7195736
## Education_LevelHigh School	0.2461248
## Education_LevelPost-Graduate	1.2457979
## Education_LevelUneducated	0.5310317
## Education_LevelUnknown	1.1791066
## Marital_StatusMarried	3.0822941
## Marital_StatusSingle	0.5893686
## Marital_StatusUnknown	0.5619667
## Income_Category\$40K - \$60K	5.4989659
## Income_Category\$60K - \$80K	4.7802932
## Income_Category\$80K - \$120K	3.0648737
## Income_CategoryLess than \$40K	4.3652086
## Income_CategoryUnknown	4.6344535
## Card_CategoryGold	2.7417825
## Card_CategoryPlatinum	2.9955012
## Card_CategorySilver	2.6497523
## Months_on_book	1.2075007
## Total_Relationship_Count	15.7457192
## Months_Inactive_12_mon	11.9818320
## Contacts_Count_12_mon	13.6722308
## Credit_Limit	8.5389789
## Total_Trans_Amt	16.5035498
## Total_Trans_Ct	28.7383614
## Avg_Utilization_Ratio	16.0475677

Of course we want to compare multiple supervised learning methods to compare the quality of different models and choose the model that is the best at predicting customer attrition. The 2nd model we use is a basic CART decision tree.

```
tree <- rpart(Attrition_Flag ~., method = "class", data = train)
printcp(tree)
plotcp(tree)
```




```
summary(tree)
```

The cp shows

```
# plot tree
plot(tree, uniform=TRUE,
      main="Classification Tree for Attrition")
text(tree, use.n=TRUE, all=TRUE, cex=.8)
```

Classification Tree for Attrition



This shows an overview of our decision tree. For instance a customer with .. and .. will have a .. percent chance of leaving the company.

We test the quality of our model by firstly predicting attrition on our training dataset.

```

#options(digits=4)
# assess the model's accuracy with train dataset by make a prediction on the train data.
Predict_model1_train <- predict(tree, train, type = "class")
#build a confusion matrix to make comparison
conMat <- confusionMatrix(as.factor(Predict_model1_train), as.factor(train$Attrition_Flag))
#show confusion matrix
conMat$table

```

```

##           Reference
## Prediction    0    1
##           0 6190  291
##           1  203  911

```

This shows the confusion matrix for our decision tree model. Our predictions on the training set seem to be doing pretty decent, but to give more insight we proceed by calculating the Accuracy, Sensitivity, Specificity and Precision of our model.

```
sensitivity(conMat$table)
```

```
## [1] 0.9682465
```

```
specificity(conMat$table)
```

```
## [1] 0.7579035
```

```
print(accuracy <- (6190+911)/(6190+911+291+203))
```

```
## [1] 0.9349572
```

```
print(precision <- 911/(911+203))
```

```
## [1] 0.8177738
```

The model looks to do a decent job, our sensitivity seems to be quite higher than our specificity, which implies that our model is better at correctly classifying clients that left than at finding true loyal customers. This could be because of ...

Now that we have constructed the model we proceed by predicting the values in the test set in order to assess the suitability of the model.

```
Predict_model1_test <- predict(tree, test, type = "class")
```

```
conMatteest <- confusionMatrix(as.factor(Predict_model1_test), as.factor(test$Attrition_Flag))
```

```
conMatteest$table
```

```
##           Reference
## Prediction      0      1
##           0 2032  110
##           1   75  315
```

This shows the confusion matrix for our decision tree model. Our predictions on the test set seem to be of similar quality as they were on the training data. Again we proceed by calculating the Accuracy, Sensitivity, Specificity and Precision of our model.

```
sensitivity(conMatteest$table)
```

```
## [1] 0.9644044
```

```
specificity(conMatteest$table)
```

```
## [1] 0.7411765
```

```
print(accuracy <- (2032+315)/(2032+315+110+75))
```

```
## [1] 0.9269352
```

```
print(precision <- 315/(315+75))
```

```
## [1] 0.8076923
```

There is not much difference between the accuracy for our model when comparing for the test and training set. The Sensitivity is slightly higher(0.01) and the specificity slightly lower(0.01). The accuracy is slightly lower than when predicting on the training set, however the difference is marginal


We construct the ROC curve to compare how well our model predicts Attrition compared to a completely random guessing strategy.

```
par(mai=c(.9,.8,.2,.2))
```

```
plot(roc(test$Attrition_Flag, as.numeric(Predict_model1_test)), print.auc=TRUE,
     col="black", lwd=1, main="ROC curve", xlab="Specificity: true negative rate", ylab="Sensitivity: t
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```



The area under the curve of ... implies that our model is doing a good job in predicting customer attrition. We calculate the variable importance for our decision tree to see which variables were decisive in classifying the customers and predicting their behaviour.

```
treevariableimportance <- varImp(tree, scale = FALSE)
print(treevariableimportance)
```

```
## Overall
## Avg_Utilization_Ratio 586.32904
## Contacts_Count_12_mon 137.52696
## Credit_Limit 205.61775
## Customer_Age 31.69373
## Gender 80.31005
## Months_Inactive_12_mon 203.82831
## Months_on_book 17.75329
## Total_Relationship_Count 413.91122
## Total_Trans_Amt 872.06459
## Total_Trans_Ct 625.78974
## Dependent_count 0.00000
## Education_Level 0.00000
## Marital_Status 0.00000
## Income_Category 0.00000
## Card_Category 0.00000
```

A 3rd type of model that we could implement is a random forest. CART decision trees are easily interpretable, however, output can be ... because of ... Therefore we implement a random forest model that uses a bagging procedure producing ... regression trees and takes the average of each regression tree to improve

the .. of the model results.

```
train$Attrition_Flag <- as.character(train$Attrition_Flag)
train$Attrition_Flag <- as.factor(train$Attrition_Flag)
rf <- randomForest(Attrition_Flag ~ ., , data = train, proximity=FALSE, importance = FALSE)
print(rf)
```

```
##
## Call:
## randomForest(formula = Attrition_Flag ~ ., data = train, , proximity = FALSE, importance = FALSE)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 3
##
##           OOB estimate of  error rate: 5.29%
## Confusion matrix:
##           0    1 class.error
## 0 6297  96  0.01501642
## 1  306 896  0.25457571
```

```
summary(rf)
```

```
##           Length Class  Mode
## call              6 -none- call
## type              1 -none- character
## predicted        7595 factor numeric
## err.rate         1500 -none- numeric
## confusion         6 -none- numeric
## votes            15190 matrix numeric
## oob.times         7595 -none- numeric
## classes           2 -none- character
## importance        15 -none- numeric
## importanceSD       0 -none- NULL
## localImportance    0 -none- NULL
## proximity          0 -none- NULL
## ntree             1 -none- numeric
## mtry              1 -none- numeric
## forest            14 -none- list
## y                 7595 factor numeric
## test              0 -none- NULL
## inbag             0 -none- NULL
## terms             3 terms  call
```

This shows the summary of our random forest model. We can see that

Just like with the 2 previous models we continue by predicting customer attrition in our training dataset to assess how well our model fits the data.

```
predrf <- predict(rf, data = "train", type = "response")
print(rftab <- table(predrf, train$Attrition_Flag))
```

```
##
## predrf      0      1
##           0 6297  306
##           1  96 896
```

```
print(accuracyrf <- (6290+914)/nrow(train))
```



```
## [1] 0.9485188
print(sensitivityrf <- 914/(914+103))
```

```
## [1] 0.8987217
print(precisionrf <- 914/(914+288))
```

```
## [1] 0.7603993
print(specificityrf <- 6290/(6290 + 288))
```

```
## [1] 0.9562177
```

The accuracy of ... implies that our model is quite accurate in its predictions. The sensitivity and specificity signal that ... whereas the precision of .. tells us ..

To verify the fit of our model we also predict customer attrition in the test dataset.

```
predrfptest <- predict(rf, newdata = test, type = "response")
print(rftabtest <- table(predrfptest, test$Attrition_Flag))
```

```
##
## predrfptest    0    1
##              0 2076   92
##              1   31  333
```

```
print(accuracyrfptest <- (2073+335)/nrow(test))
```

```
## [1] 0.9510269
print(sensitivityrfptest <- 335/(335+34))
```

```
## [1] 0.9078591
print(precisionrfptest <- 335/(335+90))
```

```
## [1] 0.7882353
print(specificityrfptest <- 2073/(2073 + 90))
```

```
## [1] 0.9583911
```


The accuracy of ... tells us that our model is quite accurate in its predictions. The sensitivity and specificity show that ... whereas the precision of .. gives us an indication of how ..

We also plot an ROC curve for our random forest model to see how well our predictions are doing compared to random guessing.

```
par(mai=c(.9,.8,.2,.2))
plot(roc(test$Attrition_Flag, as.numeric(predrfptest)), print.auc=TRUE,
     col="black", lwd=1, main="ROC curve", xlab="Specificity: true negative rate", ylab="Sensitivity: t
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```



The area under the curve of ... indicates that our model is doing a good job in predicting customer attrition.

In the following graph we compare the results of all our 3 models at the same time.

```
glm.roc <- roc(response = test$Attrition_Flag, predictor = as.numeric(predtest))

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases


rpart.roc <- roc(response = test$Attrition_Flag, predictor = as.numeric(Predict_model1_test))

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

rf.roc <- roc(response = test$Attrition_Flag, predictor = as.numeric(predrfest))

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

plot(glm.roc,      legacy.axes = TRUE, print.auc.y = 1.0, print.auc = TRUE)
plot(rpart.roc, col = "blue", add = TRUE, print.auc.y = 0.65, print.auc = TRUE)
plot(rf.roc, col = "red", add = TRUE, print.auc.y = 0.85, print.auc = TRUE)
legend("bottom", c("Random Forest", "Decision Tree", "Logistic"),
      lty = c(1,1), lwd = c(2, 2), col = c("red", "blue", "black"), cex = 0.75)
```



The graph indicates that the Logistic regression model that we constructed seems to be the best model for predicting customer attrition. The area under the curve of .. indicates that in ... of the cases we would predict correctly.

Given that the logistic regression model came out as the best model we shortly highlight the results of the logistic regression to give insights into the drivers of customer attrition. According to our regression the,, ... and where the most important predictors of customer attrition. ... with a coefficient of ... would on average increase the likelihood of attrition by ... percent.

... with a coefficient of ... would on average increase the likelihood of churning by ... percent.

... with a coefficient of ... would on average increase the likelihood of churning by ... percent.

... with a coefficient of ... would on average increase the likelihood of churning by ... percent.